



CS 240: Programming in C

Lecture 7: Arrays Memory Layout of Data

Prof. Jeff Turkstra



Announcements

- Work on Homework 3!
 - Done by Sunday!
- Feasting with Faculty tomorrow 12pm
 - Look for the sign on the door (2 private dining rooms toward back of the area)
 - You can show up to any (all!) of them!

Grades

- Homework grades are typically available the day after the assignment is due
 - No more announcements, your job to check
- Lecture quizzes same
- Takehome quizzes generally take one week to grade
 - Then released on Gradescope
 - Then in gradebook a week later
 - Regrade request deadline is based on Gradescope release

Homework 3

- What's wrong with this? (Assume we check fscanf()'s return value)

```
char buf[1024];  
fscanf(in_fp, "%[^\\n]", buf);  
if (strlen(buf) > MAX_NAME_LEN) {  
    fclose(in_fp);  
    in_fp = NULL;  
    return BAD_RECORD;  
}
```

Homework 3

- How about this?

```
#define MAX_NAME_LEN (40)
```

```
char buf[MAX_NAME_LEN];  
fscanf(in_fp, "%40[^\n]", buf);  
...
```

Homework 3

- Don't forget the NUL terminator!

```
#define MAX_NAME_LEN (40)
```

```
char buf[MAX_NAME_LEN];  
fscanf(in_fp, "%39[^\n]", buf);  
...
```

Quizzes

- Must use the template
- Must be handwritten
- Otherwise 0
 - This has been discussed previously

Debug output

- We have fairly strict file size limits on output
- If you leave a bunch of `printf()`s in your code you might hit it
- You've been warned
 - Future assignments that's a score of 0

Feasting with Faculty

- Tomorrow! 12pm!
- Earhart Private Dining Room



What about hw3.h?

```
extern char g_rental_history[MAX_RENTALS][3][MAX_BUF_LEN ];  
extern char g_vehicle_info[MAX_RENTALS][3][MAX_BUF_LEN];  
extern float g_rental_stats[MAX_RENTALS][4];  
  
extern int g_rental_count;
```

- **extern** is also a declaration
 - It tells the compiler what the variable looks like, but it does not allocate space for it!
 - You still must define it somewhere!

Notes

- Do not use variable length arrays in this class
- Some of you are copying code and concepts from things outside of course material
 - `fgetc()`, `sizeof()`, `malloc()`, etc
 - You're probably cheating.
 - You're also making your life more difficult

Reading

- In K&R:
 - Read Sections: 4.4, 6.8-6.9, A8.3-A8.4
 - ...and skim Chapter 2 (read 2.3)
- In Beej's:
 - Read Chapter 6, ignore 6.2
 - Read Chapter 14
 - Read sections 12.2-12.4

Definitions vs. declarations

- Definition: allocates storage for a variable (or function)
- Declaration: announces the properties of a variable (or function)

- What's this?

```
struct hey {  
    int    zap;  
    float  zing;  
};
```

- And this?

```
struct point {  
    int x;  
    int y;  
} var;
```

Arrays of structures

- We can create arrays of structures just as we can create arrays of anything else. E.g.:
`struct person people[4];`
- Initialization is similar to before:

```
struct person people[4] = {  
    { "Mai Elkady", "TA", {1, 2, 3, 4} },  
    { "Nan Jiang", "TA", {2, 3, 4, 5} },  
    { "Zach Bryant", "TA", {3,4,5,6} },  
    { "Julie Stevenson", "TA",  
      {4, 5, 6, 7} },  
};
```

Array of Structures

Example (page 1)

```
#include <stdio.h>
#include <string.h>
```

```
struct person {
    char name[40];
    char title[15];
    int  codes[4];
};
```

```
struct person crowd[100]; /* global! */
```

```
void print_person(struct person);
```



Array of Structures

Example (page 2)

```
int main() {  
    int index = 0;  
  
    strncpy(crowd[0].name, "Jeff", 40);  
    strncpy(crowd[0].title, "Speaker", 15);  
    crowd[0].codes[0] = 10;  
    crowd[0].codes[1] = 20;  
    crowd[0].codes[2] = 40;  
  
    strncpy(crowd[1].name, "Student", 40);  
    strncpy(crowd[1].title, "Listener", 15);  
    crowd[1].codes[0] = 1;
```



Array of Structures

Example (page 3)

```
for (index = 0; index < 100; index++) {  
    if (crowd[index].name[0] != '\0') {  
        print_person(crowd[index]);  
    }  
}  
return 0;  
}  
  
/* Assume that print_person is defined  
 * below.  
 */
```

The result...

```
$ vi ex2.c  
$ gcc -Wall -Werror -std=c99 -g -o ex2  
ex2.c  
$ ./ex2
```

```
Name:  Jeff  
Title: Speaker  
Codes: 10, 20, 0, 9
```

```
Name:  Student  
Title: Listener  
Codes: 1, 0, 0, 0
```

```
$
```



Notes about previous example

- When you define something as a global data structure, anything that is not initialized is automatically made zero
 - Sometimes this is good, sometimes not
- We only defined the first two elements of the big array
- You can check if the first character of a string is NUL by:
`if (string[0] == '\0') ...`

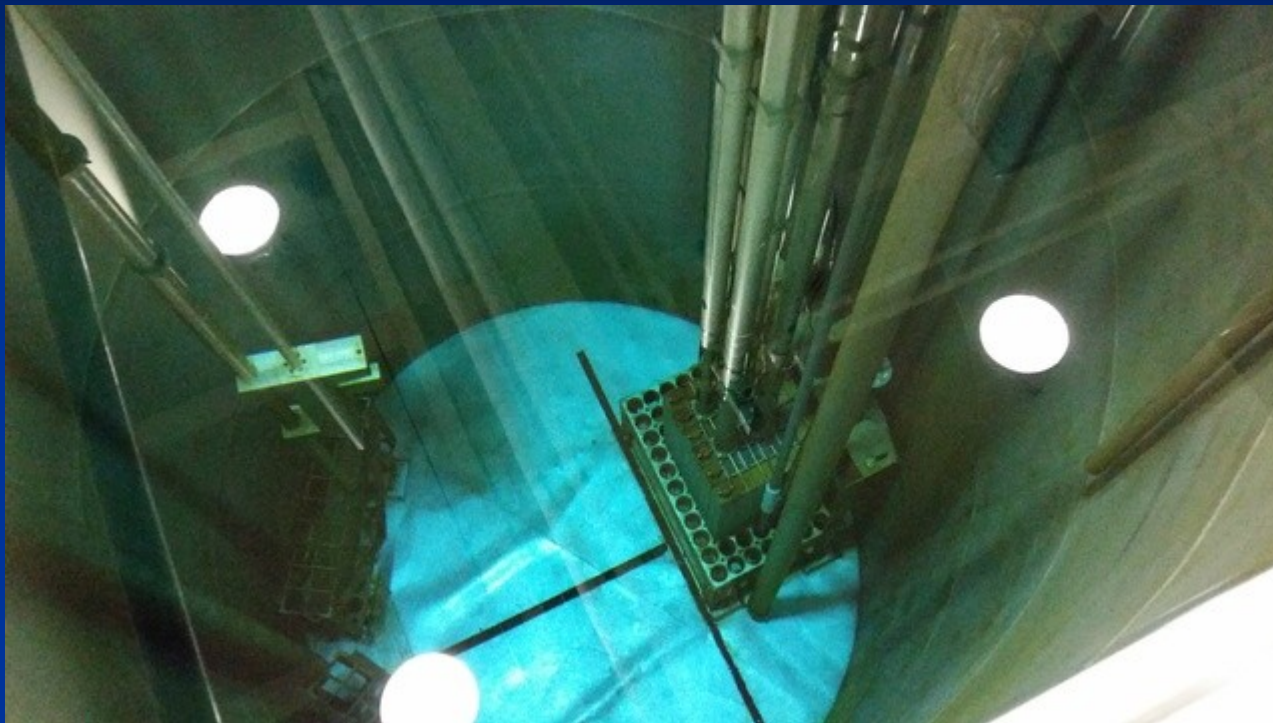
Purdue Trivia

- Purdue is home to Indiana's first and only nuclear reactor
 - Built in 1962
 - Built by Lockheed Corporation
 - Three stories beneath the Duncan Annex of EE
 - Criticality on August 30, 1962
 - Dedication September 27









Array initialization

- You can partially initialize an array! E.g.:

```
int my_numbers[200] = { 5, 5, 3, 4, 5 };
```

- Only the first five elements are explicitly initialized. The rest are set to zero
- This is true not only for global arrays but for arrays allocated inside functions

Array auto-sizing

- You can define and initialize an array without explicitly saying what its size is. E.g.:

```
int my_array[] = { 1, 1, 2, 2, 3, 3, 7 };
```

- What would the size of this array be?
- There are no zero elements at the end of the array since we're letting the compiler figure out how large it is

Arrays of structures

- Same idea...

```
struct point {  
    int x;  
    int y;  
};
```

```
int almost_pointless() {  
    struct point dots[] = { {1, 2},  
                             {3, 4} };  
    return dots[1].x;  
}
```

strncpy()

- What's wrong with this?


```
int main() {  
    char another_str[16] = "123456789012345";  
    char my_str[] = "Hello, World!";  
  
    strncpy(another_str, my_str, strlen(my_str));  
    printf("%s\n", another_string);  
  
    return 0;  
}
```

strncpy()

- Do not do this...

```
int main() {  
    char another_str[16] = "123456789012345";  
    char my_str[] = "Hello, World!";  
  
    strncpy(another_str, my_str, strlen(my_str));  
    printf("%s\n", another_string);  
  
    return 0;  
}
```

strncpy() will not NUL
terminate the string!



strncpy() fixed?


- Don't do this either...

```
int main() {  
    char another_str[16] = "123456789012345";  
    char my_str[] = "Hello, World!";  
  
    strncpy(another_str, my_str, strlen(my_str));  
    another_str[strlen(my_str)] = '\0';  
    printf("%s\n", another_string);  
  
    return 0;  
}
```

strncpy() fixed?

- Don't do this either...

```
int main() {  
    char another_str[16] = "123456789012345";  
    char my_str[] = "Hello, World!";  
  
    strncpy(another_str, my_str, strlen(my_str));  
    another_str[strlen(my_str)] = '\0';  
    printf("%s\n", another_string);  
  
    return 0;  
}
```

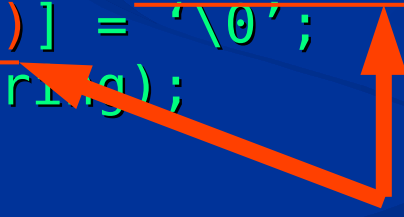


Only works because my_str happens to be smaller than another_str. What happens if my_str changes to something larger in the future?

strncpy() overflow

■ Oops...

```
int main() {  
    char another_str[16] = "123456789012345";  
    char my_str[40] = "1234567890123456789" \  
        "123456789012345678";  
  
    strncpy(another_str, my_str, strlen(my_str));  
    another_str[strlen(my_str)] = '\0';  
    printf("%s\n", another_str);  
  
    return 0;  
}
```



What's the right thing to do here?

Data layout in memory

- Everything that contains a value uses memory
- Everything that contains a value uses memory
- Everything that contains a value uses memory
- Memory space looks like a long, continuous stream of bytes



- And everything that contains a value occupies one or more bytes of memory

Variables

- When we define a variable, the compiler creates a space for it in memory somewhere. Whenever we use the name of the variable, it gets translated into that 'somewhere.'
- Some types of variables consume several bytes of memory. E.g., an 'int' is usually 4 bytes long.

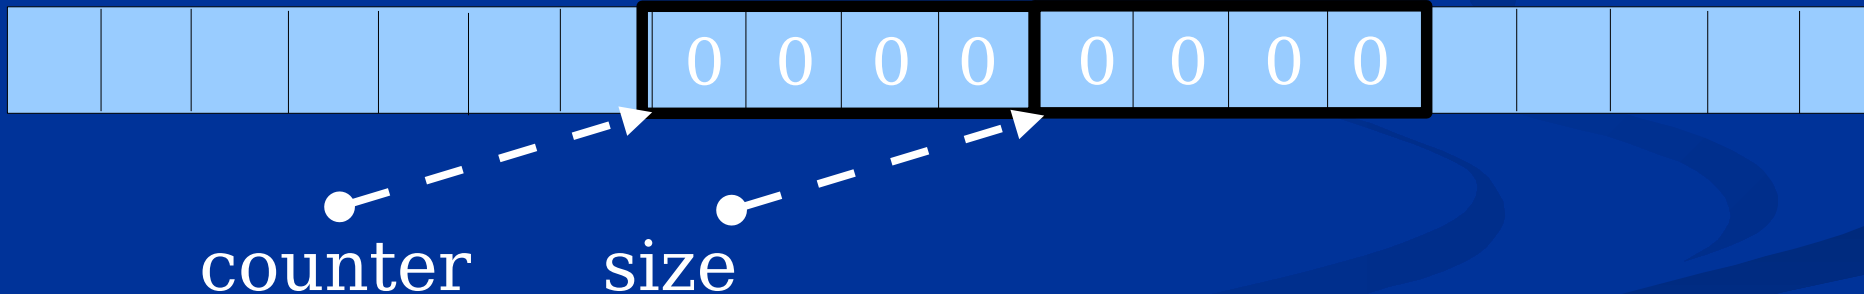
```
int my_var = 0;
```



More variables

- Variables that are defined near each other are **usually** near to each other in memory.
e.g.:

```
int counter = 0;  
float size = 0.0;
```



Arrays

- Arrays of items are guaranteed to be packed together in memory. e.g.:

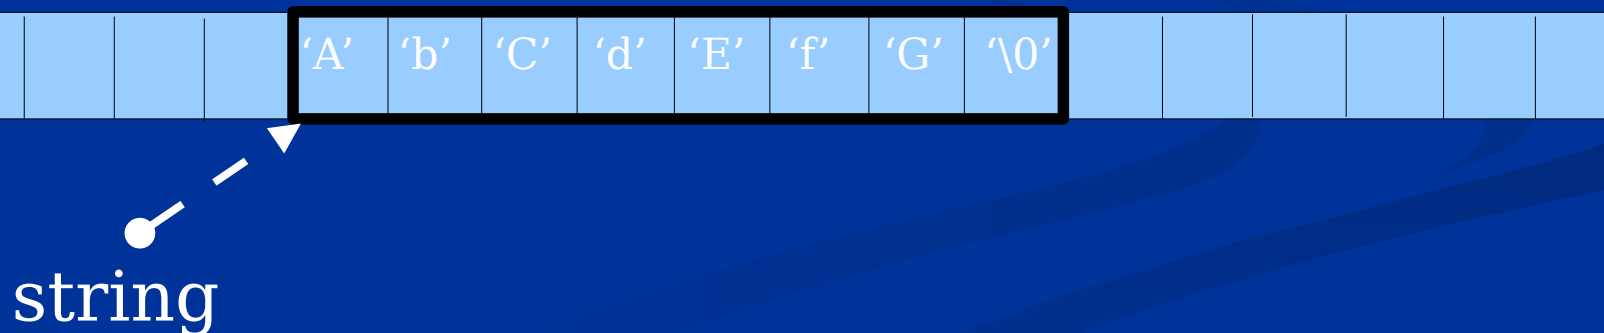
```
int array[3] = {0x11111111,  
                0x22222222,  
                0x33333333 };
```



array

Strings

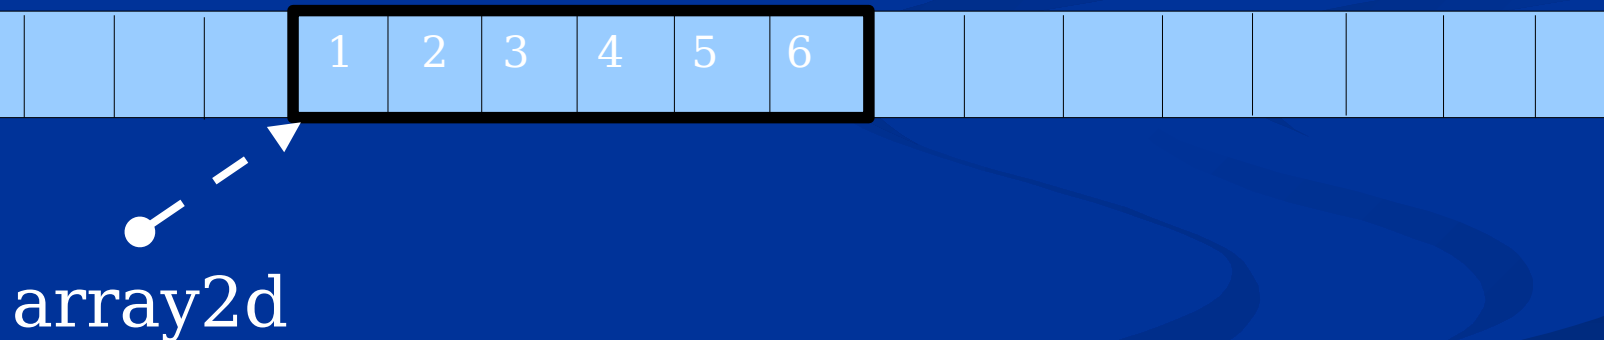
- A string in C is an array of characters
- How are these characters stored?
- All strings delimited by (") characters are said to be null-terminated (terminated by a zero byte)
- strcpy(), strcmp(), etc will search for the null. E.g.:
`char string[8] = "AbCdEfG";`



Two dimensional arrays

- How does a 2-D array get stored in memory?

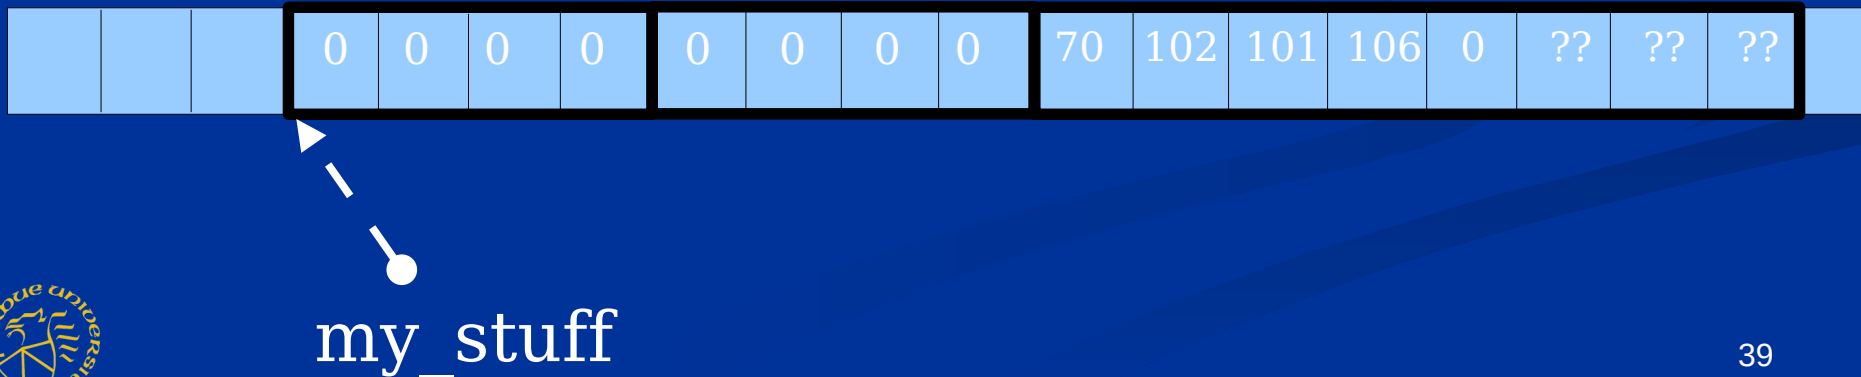
```
char array2d[2][3] = { { 1, 2, 3 },  
                      { 4, 5, 6 } };
```



Structures

- Structure members are placed in memory just like arrays...they are guaranteed to be packed next to each other.

```
struct my_stuff {  
    int i;  
    float f;  
    char c[8];  
} my_var = { 0, 0, "Ffej" };
```



How do you know the size of variables and types?

- A variable of a certain size may have a different allocated size on different machines with different compilers.
 - E.g. long X would be four bytes on x86 or Sparc but would be eight bytes long on Alpha, Sparc64, or x86_64.
- We don't want our software to misbehave when compiled on a different system.
- Fortunately, we don't have to remember what the size is...

sizeof()

- The sizeof() operator can tell us the size (number of bytes) of any:
 - Variable definition
 - Type declaration

```
int array[100];  
printf("Size of char = %d\n", sizeof(char));  
printf("Size of Array = %d\n", sizeof(array));
```

Correct strncpy()

■ :-)

```
int main() {  
    char another_str[16] = "123456789012345";  
    char my_str[40] = "1234567890123456789" \  
        "123456789012345678";  
  
    strncpy(another_str, my_str, sizeof(another_str));  
    another_str[sizeof(another_str) - 1] = '\0';  
    printf("%s\n", another_str);  
  
    return 0;  
}
```

Takehome Quiz 4

```
#include <stdio.h>
```

```
int main() {  
    char buf[11] = "Purdue";  
    int my_int = 0x8badf00d;  
    char my_char = 'X';  
    short my_short = 0xbeef;  
  
    printf("%s %d %c %hd\n", buf, my_int, my_char, my_short);  
  
    return 0;  
}
```

1. Draw the memory map as described previously
 - Remember to use `setarch -R ./your_exe` when running!
 - And run on `data.cs.purdue.edu`!
2. Are there any gaps between the space allocated for the variables?
 - If so, why might that be?

For next lecture...

- Read!
- In K&R:
 - Read Sections: 4.4, 6.8-6.9, A8.3-A8.4
 - ...and skim Chapter 2 (read 2.3)
- In Beej's:
 - Read Chapter 6, ignore 6.2
 - Read Chapter 14
 - Read sections 12.2-12.4

Boiler Up!