



CS 240: Programming in C

Lecture 25: Core Files and goto Makefiles Networking

Prof. Jeff Turkstra



Announcements

- Course Evaluations Available
- Homework 13 Extra Credit
- Hardware lecture next Monday will be mostly video
 - Includes interview with campus squirrel!
 - I'll be here for questions, though



Final Exam


- Thursday, May 8
- 10:30am – 12:30pm
- Check the seating chart
 - Available sometime before the exam
- Coding, short answer, multiple choice, true/false

Security

- You have a very, very small taste of what kind of problems can arise in terms of program security
- We've only touched the "tip of the iceberg" in terms of buffer overflows
- If you want to know more, check out:
 - <https://turkeyland.net/projects/overflow/>
 - ...or find "Smashing the Stack for Fun and Profit" using a search engine
- There are many, many other types of vulnerabilities
- If you enjoy this stuff, take a security course!

Core Files

- Does anyone know what “core” memory was?
- When your program has an unrecoverable error, the operating system saves the heap/stack memory at the exact time of the failure into a file named “core”.
- You can use the core file with the debugger



Core dump file

- `$ man 5 core`
- May have to enable it (e.g., on `data.cs.purdue.edu`)
 - `$ ulimit -c unlimited`

The Official Disclaimer with respect to “goto”

- "For a number of years I have been familiar with the observation that the quality of programmers is a decreasing function of the density of go to statements in the programs they produce. More recently I discovered why the use of the go to statement has such disastrous effects, and I became convinced that the go to statement should be abolished from all "higher level" programming languages (i.e. everything except, perhaps, plain machine code). At that time I did not attach too much importance to this discovery; I now submit my considerations for publication because in very recent discussions in which the subject turned up, I have been urged to do so."
--Edsger W. Dijkstra, March 1968, Comm. of ACM,
"Go To Statement Considered Harmful"

Why is goto bad?

- Dijkstra made the case that goto was harmful for the following reasons:
 - It prevents the compiler from being able to make “nice” computer-sciency reductions of the program
 - It makes your code unreadable
 - It is really not necessary
 - You can always rewrite to have the same functionality without goto

Why does C have a goto?

- Because...
 - The compiler doesn't have any more difficulty analyzing a program with gotos in it
 - It often makes the program clearer to read
 - It is very useful – at a certain level, at least
- Contradictions?
- More enlightened languages have even more dangerous control flow operations

What does goto look like?

- You can define labels and goto those labels...

```
int func(int x) {  
    int sum = 0;
```

```
    again:
```

```
        sum = sum + x;
```

```
        x = x - 1;
```

```
        if (x <= 0)
```

```
            goto get_out;
```

```
        else
```

```
            goto again;
```

```
    get_out:
```

```
        return sum;
```

```
}
```

How can goto make a program clearer to read?

- When you really need to ditch the control flow of your program and take drastic measures:

```
start_over:
for (int x = 0; x < 5000; x++) {
    ptr = array[x];
    while (ptr->val < level) {
        while (ptr->next != 0 && ptr->val < level) {
            if (ptr->total == 0) {
                level++;
                goto start_over;
            }
        }
        sum += ptr->total;
    }
}
```



When is goto useful?

- When it is necessary to break out of deeply nested loops (previous example)
- When you're building a state machine in software
- In general, you should still avoid using gotos unless there is a **really** good reason

Makefile

- Simple way to help organize code compilation
- Composed of rules
 - Target – usually a file to generate
 - Can be an action (“make clean”)
 - Prerequisites – used to create the target
 - Recipe – action to carry out
 - Must start with a tab!

```
gcc -o hello hello.c hellofunc.c -I.
```

Simple, hard coded

```
hello: hello.c hellofunc.c
```

```
gcc -o hello hello.c hellofunc.c -l.
```

Or...

```
CC=gcc
```

```
CFLAGS=-l.
```

```
hello: hello.o hellofunc.o
```

```
$(CC) -o hello hello.o hellofunc.o $(CFLAGS)
```


More generic

CC=gcc

CFLAGS=-I.

DEPS = hello.h

%.o: %.c \$(DEPS)

\$(CC) -c -o \$@ \$< \$(CFLAGS)

hello: hello.o hellofunc.o

gcc -o hello hello.o hellofunc.o -l.

More Variables

CC=gcc

CFLAGS=-I.

DEPS = hellomake.h

OBJ=hello.o hellofunc.o

%.o: %.c \$(DEPS)

\$(CC) -c -o \$@ \$< \$(CFLAGS)

hello: \$(OBJ)

gcc -o \$@ \$^ \$(CFLAGS)



.PHONY: clean

clean:

```
rm -f $(ODIR)/*.o *~ core $(INCDIR)/*~
```

Lot's More

- https://www.gnu.org/software/make/manual/html_node/index.html

Bubble Sort

- Sorting is a big part of computer science
- Lot's of different ways with different performance/complexity
 - More in CS 251
- Bubble sort is one approach to sorting

Bubble Sort

```
void bubble_simple(int *arr, int size) {  
    for (int k = 0; k < size; k++) {  
        for (int i = 1; i < size; i++) {  
            if (arr[i-1] > arr[i]) { // swap  
                arr[i-1] ^= arr[i];  
                arr[i] ^= arr[i-1];  
                arr[i-1] ^= arr[i];  
                print_array(arr, ARRAY_SIZE);  
            }  
        }  
    }  
}
```

Better

```
void bubble(int *arr, int size) {  
    int swapped = 1;  
  
    while (swapped) {  
        swapped = 0;  
        for (int i = 1; i < size; i++) {  
            if (arr[i-1] > arr[i]) { // swap  
                int tmp = arr[i-1];  
                arr[i-1] = arr[i];  
                arr[i] = tmp;  
                swapped = 1;  
                print_array(arr, ARRAY_SIZE);  
            }  
        }  
    }  
}
```



Networking Basics

- Given a network and a set of systems, how do we actually send data between a subset of systems?
 - First have to find the system(s)
 - Establish a route
 - Decide how we're going to communicate
 - Protocol
 - Establish a connection
 - Socket

Internet protocol

- IP is an addressing and fragmentation protocol
- Breaks communication into chunks (packets)
- Routes packets from a source to a destination
- Inherently unreliable
 - No guarantee anything will make it
 - No acknowledgments
- Different versions (IPv4 vs. IPv6)

IP address

- Each host or system has at least one IP address
 - nnn.nnn.nnn.nnn – dotted decimal notation, 4 bytes, 32 bits
 - 128.10.116.31
- Packets are sent from/to IP addresses
 - May traverse multiple routers
- Public/private
- NATs

Domain Name System

- But what about google.com?
- DNS is a **distributed system** that resolves host names to IP addresses
 - Hierarchical
 - Root servers
 - Many authoritative servers
 - Name resolution can involve multiple queries
 - Caching servers
- Reverse lookups

Transmission Control Protocol

- Remember, IP is unreliable
- TCP builds on IP to create a reliable network connection (often referred to as TCP/IP)
- Supports acknowledgments and retransmission
- Hosts identified by IP address and a port number

Clients and servers

- A **server** is a process that waits for a connection
- A **client** is a process that connects to a server
- Processes can be both!
- Machines can have multiple servers and clients running at once
- Once connected, clients and servers read and write data from/to each other much like a file

Networking Basics

- Given a network and a set of systems, how do we actually send data between a subset of systems?
 - First have to find the system(s)
 - Establish a route
 - Decide how we're going to communicate
 - Protocol
 - Establish a connection
 - Socket

Internet protocol

- IP is an addressing and fragmentation protocol
- Breaks communication into chunks (packets)
- Routes packets from a source to a destination
- Inherently unreliable
 - No guarantee anything will make it
 - No acknowledgments
- Different versions (IPv4 vs. IPv6)

IP address

- Each host or system has at least one IP address
 - nnn.nnn.nnn.nnn – dotted decimal notation, 4 bytes, 32 bits
 - 128.10.116.31
- Packets are sent from/to IP addresses
 - May traverse multiple routers
- Public/private
- NATs

Domain Name System

- But what about google.com?
- DNS is a **distributed system** that resolves host names to IP addresses
 - Hierarchical
 - Root servers
 - Many authoritative servers
 - Name resolution can involve multiple queries
 - Caching servers
- Reverse lookups

Transmission Control Protocol

- Remember, IP is unreliable
- TCP builds on IP to create a reliable network connection (often referred to as TCP/IP)
- Supports acknowledgments and retransmission
- Hosts identified by IP address and a port number

Clients and servers

- A **server** is a process that waits for a connection
- A **client** is a process that connects to a server
- Processes can be both!
- Machines can have multiple servers and clients running at once
- Once connected, clients and servers read and write data from/to each other much like a file

Network sockets

- Clients and servers communicate via sockets
- A socket is an endpoint for sending and receiving data
- A socket address consists of the IP address and port number
 - At least for TCP
- Port number is 16-bits (0-65535)
 - Ports < 1024 are privileged
- 128.10.116.31:80

Server

- `man 7 ip`
- Steps for listening...
 - Create a `socket()`
 - `bind()` that socket to an address and port
 - `listen()` for a connection
 - `accept()` the connection
 - [communicate – read, write, recv, send]
 - `close()` the connection

socket()

- Create an endpoint for communication

```
int socket(int domain, int type, int protocol);
```

bind()

- “assigning a name to a socket”

```
bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

- **socket()** gives us a fd
- **bind()** assigns an “address” to the socket

```
struct sockaddr_in {  
    sa_family_t    sin_family; /* address family: AF_INET */  
    in_port_t      sin_port;   /* port in network byte order */  
    struct in_addr sin_addr;    /* internet address */  
};
```

listen()

- Marks the socket as a passive socket
 - Accepts incoming connections

`int listen(int sockfd, int backlog);`

- `backlog` is the maximum length for the pending connections queue
 - Max number of waiting connections

accept()

```
int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```

- Removes first connection request from pending connections queue
 - Must be called on a **listen()**ing socket
 - `addr` is filled with peer information
- Creates a **new** connected socket
 - This socket does not listen
 - Original socket is left alone

recv() and send()

- **recv()** is **read()** when **flags = 0**
 - Can behave differently depending on the flag(s)
- **send()** is **write()** when **flags = 0**
 - Can behave differently depending on the flag(s)

close()

- You should always check the return value of close :-)

Client

- Steps for connecting...
 - Create a **socket()**
 - Optionally **bind()**
 - For a specific source port
 - **connect()** to an address:port
 - [communicate – read, write, recv, send]
 - **close()** the connection

connect()

```
int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

- Connects sockfd to address addr

server.c

- Simple echo server
- Can connect using **telnet**
 - Consider writing your own client – it's not that hard!

Boiler Up!