

CS 240: Programming in C

Lecture 7: Arrays, Memory Layout of Data

Announcements

- Homework 2 due tonight!
- Work on homework 3

Homework 3

hw3.h

```
extern int g_mission_count;  
extern char g_astronaut_and_mission[MAX_MISSIONS][3][MAX_NAME_LEN];  
extern char g_equipment[MAX_MISSIONS][2][MAX_NAME_LEN];  
extern float g_op_hours[MAX_MISSIONS][2];  
extern char g_mission_dates[MAX_MISSIONS][2][MAX_DATE_LEN];  
extern int g_experiments[MAX_MISSIONS][2];
```

- *extern* is a *declaration*
 - It tells the compiler what the variable looks like, but it does not allocate space for it!
 - You still must define it somewhere!

Arrays of structures

- We can create arrays of structures just as we can create arrays of anything else

```
struct person people[4];
```

- Initialization is similar to before:

```
struct person people[4] = {  
    { "Kirk",  "Captain",          {10, 20, 25, 9} },  
    { "Spock", "First Officer",    {-1,  2,  5, 8} },  
    { "McCoy", "Chief Medical Officer", {94, 33,  2, 8} },  
    { "Scott", "Chief Engineer",    { 7,  7,  2, 8} },  
};
```

Arrays of structures example

```
#include <stdio.h>
#include <string.h>

struct person {
    char name[40];
    char title[25];
    int  codes[4];
};

struct person crew[200]; /* global! */

void print_person(struct person);
```

Arrays of structures example

```
int main() {  
    int index = 0;  
  
    strncpy(crew[0].name, "Kirk", 40);  
    strncpy(crew[0].title, "Captain", 25);  
    crew[0].codes[0] = 10;  
    crew[0].codes[1] = 20;  
    crew[0].codes[2] = 40;  
  
    strncpy(crew[1].name, "Ensign", 40);  
    strncpy(crew[1].title, "Redshirt", 25);  
    crew[1].codes[0] = 1;  
}
```

Arrays of structures example

```
int index;  
for (index = 0; index < 200; index++) {  
    if (crew[index].name[0] != '\0') {  
        print_person(crew[index]);  
    }  
}  
return 0;  
}
```

```
/* Assume that print_person is defined  
 * below.  
 */
```

Arrays of structures example

```
$ vi ex2.c
$ gcc -Wall -Werror -std=c99 -g -o ex2 ex2.c
$ ./ex2
Name:  Kirk
Title: Captain
Codes: 10, 20, 40, 0

Name:  Ensign
Title: Redshirt
Codes: 1, 0, 0, 0

$
```


Notes about previous example

- When you define something as a global data structure, anything that is not initialized is automatically made zero
 - Sometimes this is good, sometimes not
- We only defined the first two elements of the big array
- You can check if the first character of a string is NUL by:

```
if (string[0] == '\0') ...
```

Array initialization

- You can partially initialize an array!

```
int my_numbers[200] = { 5, 5, 3, 4, 5 };
```

- Only the first five elements are explicitly initialized. The rest are set to zero
- This is true not only for global arrays but for arrays allocated inside functions as well

Array auto-sizing

- You can define and initialize an array without explicitly saying what its size is

```
int my_array[] = { 1, 7, 0, 1 };
```

- This array would be size 4
- There are no zero elements at the end of the array since we're letting the compiler figure out how large it is

Arrays of structures

- Same idea...

```
struct point {  
    int x;  
    int y;  
};  
  
int almost_pointless() {  
    struct point dots[] = { {1, 2},  
                             {3, 4} };  
    return dots[1].x;  
}
```

strncpy()

- What's wrong with this?

```
int main() {  
    char another_str[16] = "123456789012345";  
    char my_str[] = "Hello, world!";  
  
    strncpy(another_str, my_str, strlen(my_str));  
    printf("%s\n", another_str);  
  
    return 0;  
}
```

strncpy()

- What's wrong with this?

```
int main() {  
    char another_str[16] = "123456789012345";  
    char my_str[] = "Hello, world!";  
  
    strncpy(another_str, my_str, strlen(my_str));  
    printf("%s\n", another_str);  
  
    return 0;  
}
```

↑
strncpy() will not NUL
terminate the string!

strncpy()

- Does this fix it?

```
int main() {  
    char another_str[16] = "123456789012345";  
    char my_str[] = "Hello, world!";  
  
    strncpy(another_str, my_str, strlen(my_str));  
    another_str[strlen(my_str)] = '\0';  
    printf("%s\n", another_str);  
  
    return 0;  
}
```

strncpy() overflow

- No! What if my_str is longer than another_str?

```
int main() {  
    char another_str[16] = "123456789012345";  
    char my_str[40] = "12345678901234567890" \  
                    "1234567890123456789";  
  
    strncpy(another_str, my_str, strlen(my_str));  
    another_str[strlen(my_str)] = '\0';  
    printf("%s\n", another_str);  
  
    return 0;  
}
```


strncpy()

- What about now?

```
int main() {  
    char another_str[16] = "123456789012345";  
    char my_str[40] = "12345678901234567890" \  
        "1234567890123456789";  
  
    strncpy(another_str, my_str, strlen(another_str));  
    another_str[strlen(another_str)] = '\0';  
    printf("%s\n", another_str);  
  
    return 0;  
}
```


strncpy()

- No! What if another_str is shorter than its buffer?

```
int main() {  
    char another_str[16] = "1";  
    char my_str[40] = "12345678901234567890" \  
                    "1234567890123456789";  
  
    strncpy(another_str, my_str, strlen(another_str));  
    another_str[strlen(another_str)] = '\0';  
    printf("%s\n", another_str);  
  
    return 0;  
}
```


Data layout in memory

- Everything that contains a value uses memory
- Memory space looks like a long, continuous stream of bytes
- And everything that contains a value occupies one or more bytes of memory



Variables


- When we define a variable, the compiler creates a space for it in memory somewhere
- Whenever we use the name of the variable, it gets translated into that “somewhere”
- Some types of variables consume several bytes of memory
 - An int is usually 4 bytes long



Variables

- Variables that are defined near each other are usually near to each other in memory


```
int counter = 0;  
float size = 0.0;
```



Arrays

- Arrays of items are **guaranteed** to be packed together in memory

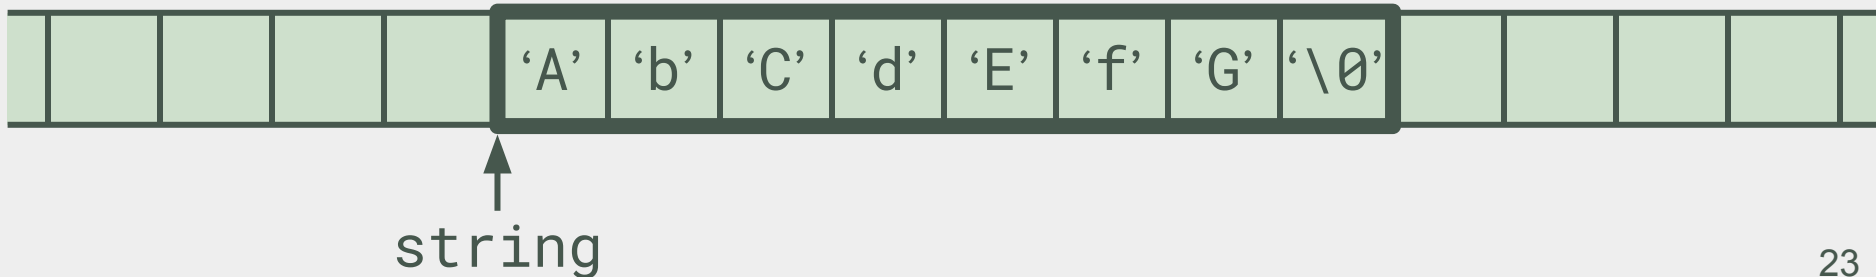
```
int array[3] = { 0x11111111,  
                0x22222222,  
                0x33333333 };
```



Strings

- A string in C is an **array** of characters
- All strings delimited by “ are said to be null-terminated
 - Terminated by a zero byte
- strcpy(), strcmp(), etc. will search for the null


```
char string[8] = "AbCdEfG";
```



Two dimensional arrays

- How does a 2-D array get stored in memory?

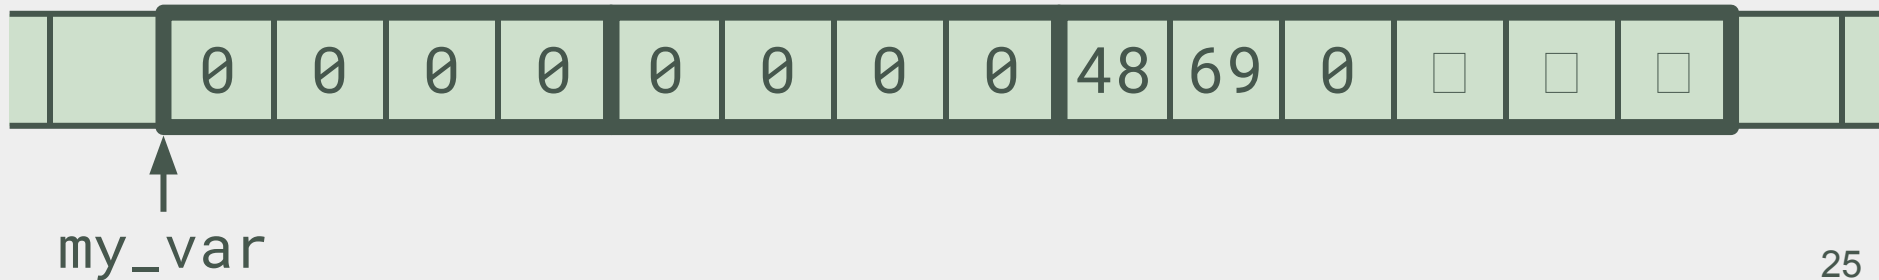
```
char array2d[2][3] = { {1, 2, 3},  
                       {4, 5, 6} };
```



Structures

- Structure members are placed in memory just like arrays
- They are guaranteed to be packed next to each other

```
struct my_stuff {  
    int i;  
    float f;  
    char c[6];  
} my_var = { 0, 0, "Hi" };
```



Variable sizes and types

- How do you know the size and type of a variable?
- It may have a different allocated size on different machines with different compilers
 - A long would be 4 bytes on x86, but would be 8 bytes on x86_64
- We don't want our software to misbehave when compiled on a different system.
- Fortunately, we don't have to remember what the size is

sizeof()

- The sizeof() operator can tell us the size (number of bytes) of any:
 - Variable definition
 - Type declaration

```
int array[100];  
printf("Size of char  = %d\n", sizeof(char));  
printf("Size of array = %d\n", sizeof(array));
```

Correct *strncpy()*

- Use `sizeof()` to get the size of the buffer!

```
int main() {  
    char another_str[16] = "1";  
    char my_str[40] = "12345678901234567890" \  
                    "1234567890123456789";  
  
    strncpy(another_str, my_str, sizeof(another_str));  
    another_str[sizeof(another_str) - 1] = '\\0';  
    printf("%s\\n", another_str);  
  
    return 0;  
}
```

What is the size of this struct?

```
struct strange {  
    int x;  /* four bytes */  
    int y;  /* four bytes */  
    int z;  /* four bytes */  
    char c; /* one byte */  
};  
  
int main() {  
    printf("size = %d\n", sizeof(struct strange));  
    return 0;  
}
```

```
$ ./strange  
size = 16
```

Right...

- The size of the previous structure is 16 bytes
- On most modern computers, an integer must reside on an even boundary if it is to be efficiently accessed

```
int my_int;  
assert(&my_int % 4 == 0)  
  
short my_short;  
assert(&my_short % 2 == 0)
```

Padding

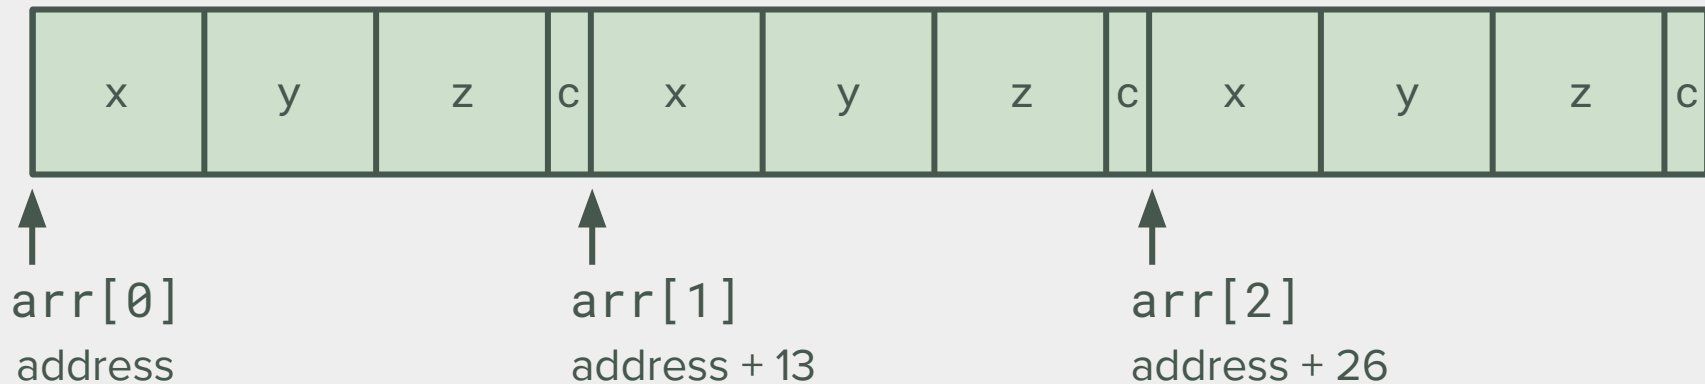
- Structures are often padded so that data elements occur at the correct offset
 - E.g., ints must be 4-byte aligned, longs must be 8-byte aligned, etc.
- Some architectures cannot handle unaligned accesses
- For others, they are very slow

If a structure is not padded

- If the structure was not padded, an array of these structures would look like this:

```
struct strange arr[3];
```

Incorrect!




- If address is a proper location for an int, address + 13 isn't

When a structure is padded

- When an odd-sized array is created, it is padded to align all of its fields properly

```
struct strange arr[3];
```

Correct!



- Now all of its integers are on a proper boundary

When a structure is padded

- You can't (shouldn't) access the pad space
- Note that padding may be added at several places in the structure

How to create inefficient structs


- Here's a structure that uses space inefficiently

```
struct bad {  
    char c1;  
    int i1;  
    char c2;  
    int i2;  
    char c3;  
    int i3;  
    char c4;  
}; /* size = 28 */
```


- How big would you say it is?

How to create inefficient structs

- All values must be properly aligned...



- And must remain aligned in an array...



How to create efficient structs

- We can simply reorder the variables to reduce padding

```
struct bad {  
    char c1;  
    int i1;  
    char c2;  
    int i2;  
    char c3;  
    int i3;  
    char c4;  
}; /* size = 28 */
```

```
struct not_so_bad {  
    int i1;  
    int i2;  
    int i3;  
    char c1;  
    char c2;  
    char c3;  
    char c4;  
}; /* size = 16 */
```

Structure alignment rule of thumb

- When creating a structure, order the fields top to bottom by their relative size
 - double
 - long long
 - pointer
 - long
 - int/float
 - short
 - char
- Doing so will result in padding being added only to the end of the structure - if at all

Takehome Quiz #2

- Due 9/12 at 2:20 pm
- Use the template!
 - Solutions not using the template will not get full credit
- Handwritten only!
 - Or no credit

Takehome Quiz #2

Run the program on the next slide, and answer these questions:

1. Draw the memory map as in slides 19-25

- a. Use `setarch -R ./your_program` to run your program
- b. Hint: use “%p” with `printf` to print the address of a variable, e.g.:
- c. And run on `data.cs.purdue.edu`

```
printf("%p\n", &var);
```

2. Are there any gaps between space allocated for the variables? If so, why?

- a. Hint: padding only applies to structures

3. The order of variables in memory may be unexpected. Can you explain it?

Takehome Quiz #2

```
#include <stdio.h>

int main() {
    char buf[6] = "Hi!";
    int my_int = 0xbeefbeef;
    struct a_struct {
        long l;
        char c;
    } my_struct = { 1701, 'Z' };
    short my_short = 0xf00d;

    printf("%s %d %ld %c %hd\n", buf, my_int, my_struct.l,
                                                my_struct.c, my_short);

    return 0;
}
```

For next lecture

- Read K&R 2.3, 4.4, 6.8-6.9, A8.3-A8.4
 - ...and skim K&R 2
- Practice the examples!!

Slides

- Slides are heavily based on Prof. Turkstra's material from previous semesters.