

# *CS 240: Programming in C*

## Lecture 8: Binary File I/O, Unions, Enums, Bitwise Operators

# *Announcements*

- Homework 4 is out
  - Try out hw4\_view

# Homework 2 histogram

291 scores total...

```
100+: (0)
100: ===== (179)
 90: == (10)
 80: == (9)
 70: = (6)
 60: == (9)
 50: = (3)
 40: = (1)
 30: == (14)
 20: === (17)
 10: = (8)
  0: ===== (35)
```

# Announcements

- Midterm 1 is next week!
  - 9/24 at 8:00 pm
  - CL50 224 and MATH 175
  - Seating chart will be released soon
- Practice questions and practice midterm on website
- You will have to handwrite code

# Binary file I/O

- Given an open FILE pointer, we can use fread() and fwrite() to read or write “raw” memory items to or from a file

```
fwrite(void *ptr, int size, int num, FILE *fp);  
fread(void *ptr, int size, int num, FILE *fp);
```

- This allows you to “dump” a data structure directly into a **binary** format file

# Example

```
#include <stdio.h>

struct xx {
    int x;
    int y;
};

int main() {
    struct xx try = { -1, -1 };
    FILE *fp = fopen("input.file", "rb");
    int status = fread(&try, sizeof(struct xx), 1, fp);
    printf("Read values (%d, %d) with return %d\n",
           try.x, try.y, status);
    fclose(fp);
    return 0;
}
```

# *Return values*

- Both fread() and fwrite() return the number of **items** that were read or written
- On error, they return a short item count (or zero)

# Uses of fread()

- Recall the prototype:

```
fread(void *ptr, int size, int num, FILE *fp);
```

```
fread(&try, sizeof(struct xx), 1, fp);
```

- The “void \*” means we can pass a pointer to “anything”
- What value should this call to fread() return?
- How many bytes are read by this operation?
- How would we read a whole file full of the structures?
- Is there any data format checking with this?



# ***fwrite()** example*

```
#include <stdio.h>

int main() {
    struct xx try = { 1, 2 };
    FILE *fp = fopen("xx.out", "wb");
    if (fp == NULL) {
        return -1;
    }

    int ret = fwrite(&try, sizeof(struct xx), 1, fp);
    printf("ret: %d\n", ret);

    fclose(fp);
    fp = NULL;
    return 0;
}
```

# ***fwrite()ing multiple structs***

```
#include <stdio.h>

int main() {
    struct xx xxs[20];
    for (int i = 0; i < 20; i++) {
        xxs[i].x = i + 1;
        xxs[i].y = -(i + 1);
    }

    FILE *fp = fopen("xxs.out", "wb");
    int ret = fwrite(xxs, sizeof(struct xx), 20, fp);
    printf("ret: %d\n\n", ret);

    fclose(fp);
    return 0;
}
```

# ***fwrite()**ing multiple structs*

```
$ vi ex1.c  
$ gcc -o ex1 ex1.c  
$ ./ex1  
ret: 20  
  
$ cat xxs.out  
  
$
```

# *fwrite()ing multiple structs*

```
$ hexdump -C xxs.out
```

00000000	01	00	00	00	ff	ff	ff	ff	02	00	00	00	fe	ff	ff	ff	.....
00000010	03	00	00	00	fd	ff	ff	ff	04	00	00	00	fc	ff	ff	ff	.....
00000020	05	00	00	00	fb	ff	ff	ff	06	00	00	00	fa	ff	ff	ff	.....
00000030	07	00	00	00	f9	ff	ff	ff	08	00	00	00	f8	ff	ff	ff	.....
00000040	09	00	00	00	f7	ff	ff	ff	0a	00	00	00	f6	ff	ff	ff	.....
00000050	0b	00	00	00	f5	ff	ff	ff	0c	00	00	00	f4	ff	ff	ff	.....
00000060	0d	00	00	00	f3	ff	ff	ff	0e	00	00	00	f2	ff	ff	ff	.....
00000070	0f	00	00	00	f1	ff	ff	ff	10	00	00	00	f0	ff	ff	ff	.....
00000080	11	00	00	00	ef	ff	ff	ff	12	00	00	00	ee	ff	ff	ff	.....
00000090	13	00	00	00	ed	ff	ff	ff	14	00	00	00	ec	ff	ff	ff	.....
000000a0																	

File offset

Contents

ASCII  
representation

# *fwrite()ing multiple structs*

- What's the difference between:

```
int ret = fwrite(xxs, sizeof(struct xx), 20, fp);
```

```
int ret = fwrite(xxs, sizeof(xxs), 1, fp);
```

- What will ret be in both cases? Assuming no errors
- Will the output files be different?

# *fread()ing multiple structs*

- We can do the same thing with fread()

```
int ret = fread(xxs, sizeof(struct xx), 20, fp);
```

```
int ret = fread(xxs, sizeof(xxs), 1, fp);
```

# Summary of fread()/fwrite()

- Moves a “memory image” to or from a file
- The file is **not portable**
- Different systems have different formats for integers, floats, etc.
  - Endianness
- No data type checking

# Endianness

- Recall we wrote the integer value 1 to a file...

```
$ hexdump -C xxs.out
00000000 01 00 00 00 ff ff ff ff 02 00 00 00 fe ff ff ff |.....|
```


- Endianness is the order of **bytes** in a word or multi-byte value
  - It does not impact **bit** ordering for individual bytes!
- Two schemes:
  - Big-endian: most significant byte first (lowest address)
  - Little-endian: least significant byte first



# Example

- Consider the integer value 305419896
- In hexadecimal:
  - 0x12345678
- Each pair of hexadecimal values corresponds to 8 bits or 1 byte

# When stored in memory...



- Each box is 1 byte. We can look at it in binary too:
  - $0x12 = 0b00010010$
  - $0x34 = 0b00110100$
  - $0x56 = 0b01010110$
  - $0x78 = 0b01111000$
- $0x12345678 = 0b00010010\ 00110100\ 01010110\ 01111000$
- $0x78563412 = 0b01111000\ 01010110\ 00110100\ 00010010$

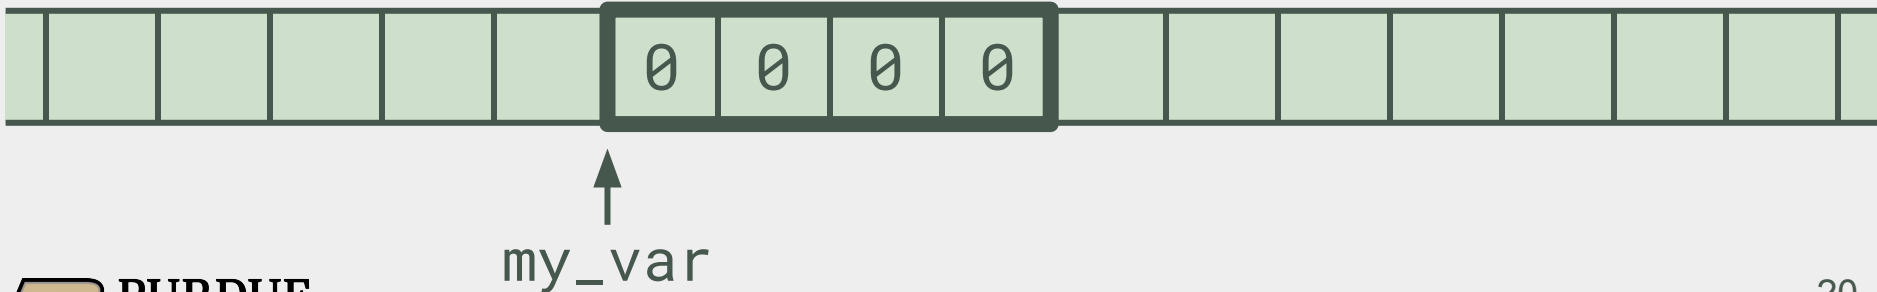
# *Binary file portability*

- Most machines use little endian
- Some file formats specify endianness
  - e.g., JPEG is always stored in Big endian
  - Little endian machines must convert the values
- Some file formats test endianness
  - E.g., read a known value from the file, check the first byte

# union

- Similar to structs, but internal elements overlap

```
union my_union {  
    int i;  
    float f;  
    char c;  
} my_var;
```



# *Union example*


```
union my_union my_var;  
my_var.i = 86593;  
  
printf("%c\n", my_var.c);
```

A

# Union example

```
union my_union my_var;  
my_var.i = 86593;  
  
printf("%c\n", my_var.c);
```

A



↑  
my\_var

# Initialization

```
union my_union {  
    int i;  
    float f;  
};  
union my_union my_var = { 5 };
```

- Assumes you are initializing the first field!
- C99 has designated union initializers:

```
union my_union my_var = { .f = 5.0 };
```

# Why unions?

- When you really need to save space in your program and you know that some data will be one of two types
- Deep operating system hacking
  - Peripheral I/O manipulation
- Format conversion
- If you need it, you'll know
  - Don't use it in this class



# *enum*

- Attaches labels to values

```
enum color {  
    RED,  
    GREEN,  
    BLUE  
};  
enum color my_hue = GREEN;
```

# *enum example*

```
#include <stdio.h>

enum color {RED, GREEN, BLUE};
int main() {
    enum color my_hue = GREEN;

    switch (my_hue) {
        case RED:
        case GREEN:
            printf("Red or Green.\n");
            break;
        case BLUE:
            printf("Blue.\n");
            break;
    }
    return 0;
}
```

## *enums can also have values*

- You can assign exact values to the enum declaration's members

```
enum british_transport {  
    LAND=1,  
    SEA=2,  
    AIR=3,  
    SUBMARINE=2,  
    FLYING_SAUCER=400  
};
```

# *Use of that enum*

```
#include <stdio.h>

int main() {
    enum british_transport craft = AIR;
    printf("Value of craft is %d\n", craft);
    return 0;
}
```

# Bitwise operators

- You regularly use logical operators:
  - `||`, `&&` in compound if statements

- What does this mean?

```
if (x) printf("x = %d\n", x);
```

- And this?

```
if (x && y) printf("x = %d\n", x);
```

- There are also bitwise operators: `|` & `&`
- What does this mean?

```
if (x & y) printf("x = %d\n", x);
```

# *The difference between && & &*

- Logical operators check whether the quantities are zero or non-zero

```
if (x && y) { ... }
```

...really means:

```
if ((x != 0) && (y != 0)) { ... }
```

- The result of && is either 1 or 0
- Use logical operators to make a **yes/no** decision

# The difference between && & &

- Bitwise operators work on all of the bits

```
char x = 5;  /* binary 00000101 */  
char y = 6;  /* binary 00000110 */  
char z = 0;  /* binary 00000000 */  
  
z = x & y;    /* result 00000100 */
```

- There are also OR (|), XOR (^) and NOT (~) operators
- Use bitwise operators when you want to work on the **bits** of a quantity

# Truth tables

AND

X	Y	O
0	0	0
0	1	0
1	0	0
1	1	1

OR

X	Y	O
0	0	0
0	1	1
1	0	1
1	1	1

XOR

X	Y	O
0	0	0
0	1	1
1	0	1
1	1	0

NOT

X	O
0	1
1	0



# *We also have shift operators*

- You can take a bunch of bits and shift them left or right

```
char x = 10;  /* binary 00001010 */  
char y = 0;  
  
y = x << 3;   /* result 01010000 */
```

- Note that every shift left is equivalent to a multiplication by two
- Above statement is equivalent to:

```
y = x * 23
```

## *Example: cut a range of bits*

- Suppose we want to write a function that accepts a 32-bit integer and pulls a range of bits from somewhere in the middle:

```
unsigned int bit_range(unsigned int num,  
                      unsigned int bits,  
                      unsigned int offset) {  
    return ((num >> offset) & ((1 << bits) - 1));  
}
```

- You'll have to stare at that for a while to understand it...

# Bit setting / clearing

- Use bitwise operators to set/clear bits in numbers...

```
int color = 44;  /* binary 00101100 */
int blue = 7;    /* binary 00000111 */

printf("Color with all blue is %d\n",
       color | blue); /* 00101111 */
printf("Color with no blue is %d\n",
       color & ~blue); /* 00101000 */
```

# Bit checking

- How can we determine if a specified bit is set?
  - i.e., set to 1

```
char bits = 44;  /* binary 00101100 */
char mask = 8;   /* binary 00001000 */

if ((bits & mask) == mask) {
    printf("The bit is set\n");
}
else {
    printf("The bit is cleared\n");
}
```

# *For next lecture*

- Read K&R Chapter 5
  - ...and/or Beej Chapter 7
  - Probably repeatedly
- Understand the operators & and \*

# *Slides*

- Slides are heavily based on Prof. Turkstra's material from previous semesters.